


# CS6700 : Tutorial 1 - Multi-Arm Bandits

Girish Madhavan V : ME20B072



Goal: Analysis 3 types of sampling strategy in a MAB

## Import dependencies

```
# !pip install seaborn

import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from typing import NamedTuple, List
```

## Gaussian Bandit Environment

```
class GaussianArm(NamedTuple):
    mean: float
    std: float

class Env:
    def __init__(self, num_arms: int, mean_reward_range: tuple, std:
float):
        """
        num_arms: number of bandit arms
```

```

    mean_reward_range: mean reward of an arm should lie between the
    given range
    std: standard deviation of the reward for each arm
    """
    self.num_arms = num_arms
    self.arms = self.create_arms(num_arms, mean_reward_range, std)

    def create_arms(self, n: int, mean_reward_range: tuple, std: float)
    -> dict:
        low_rwd, high_rwd = mean_reward_range
        # creates "n" number of mean reward for each arm
        means = np.random.uniform(low=low_rwd, high=high_rwd, size=(n,))
        arms = {id: GaussianArm(mu, std) for id, mu in enumerate(means)}
        return arms

    @property
    def arm_ids(self):
        return list(self.arms.keys())

    def step(self, arm_id: int) -> float:
        arm = self.arms[arm_id]
        return np.random.normal(arm.mean, arm.std)    # Reward

    def get_best_arm_and_expected_reward(self):
        best_arm_id = max(self.arms, key=lambda x: self.arms[x].mean)
        return best_arm_id, self.arms[best_arm_id].mean

    def get_avg_arm_reward(self):
        arm_mean_rewards = [v.mean for v in self.arms.values()]
        return np.mean(arm_mean_rewards)

    def plot_arms_reward_distribution(self, num_samples=1000):
        """
        This function is only used to visualize the arm's distribution.
        """
        fig, ax = plt.subplots(1, 1, sharex=False, sharey=False,
        figsize=(9, 5))
        colors = sns.color_palette("hls", self.num_arms)
        for i, arm_id in enumerate(self.arm_ids):
            reward_samples = [self.step(arm_id) for _ in range(num_samples)]
            sns.histplot(reward_samples, ax=ax, stat="density", kde=True,
            bins=100, color=colors[i], label=f'arm_{arm_id}')
            ax.legend()
        plt.show()

```

## Policy

```

class BasePolicy:
    @property
    def name(self):

```

```

    return 'base_policy'

def reset(self):
    """
    This function resets the internal variable.
    """
    pass

def update_arm(self, *args):
    """
    This function keep track of the estimates
    that we may want to update during training.
    """
    pass

def select_arm(self) -> int:
    """
    It returns arm_id
    """
    raise Exception("Not Implemented")

```

## Random Policy

```

class RandomPolicy(BasePolicy):
    def __init__(self, arm_ids: List[int]):
        self.arm_ids = arm_ids

    @property
    def name(self):
        return 'random'

    def reset(self) -> None:
        """No use."""
        pass

    def update_arm(self, *args) -> None:
        """No use."""
        pass

    def select_arm(self) -> int:
        return np.random.choice(self.arm_ids)

class EpGreedyPolicy(BasePolicy):
    def __init__(self, epsilon: float, arm_ids: List[int]):
        self.epsilon = epsilon
        self.arm_ids = arm_ids
        self.Q = {id: 0 for id in self.arm_ids}
        self.num_pulls_per_arm = {id: 0 for id in self.arm_ids}

    @property

```

```

def name(self):
    return f'ep-greedy ep:{self.epsilon}'

def reset(self) -> None:
    self.Q = {id: 0 for id in self.arm_ids}
    self.num_pulls_per_arm = {id: 0 for id in self.arm_ids}

def update_arm(self, arm_id: int, arm_reward: float) -> None:
    self.num_pulls_per_arm[arm_id] += 1
    old_q = self.Q[arm_id]
    n = self.num_pulls_per_arm[arm_id]
    new_q = old_q + (arm_reward - old_q) / n
    self.Q[arm_id] = new_q

def select_arm(self) -> int:
    if np.random.rand() < self.epsilon:
        return np.random.choice(self.arm_ids)
    else:
        best_arm = max(self.Q, key=self.Q.get)
        return best_arm

class SoftmaxPolicy(BasePolicy):
    def __init__(self, tau, arm_ids):
        self.tau = tau
        self.arm_ids = arm_ids
        self.Q = {id: 0 for id in self.arm_ids}
        self.num_pulls_per_arm = {id: 0 for id in self.arm_ids}

    @property
    def name(self):
        return f'softmax tau:{self.tau}'

    def reset(self):
        self.Q = {id: 0 for id in self.arm_ids}
        self.num_pulls_per_arm = {id: 0 for id in self.arm_ids}

    def update_arm(self, arm_id: int, arm_reward: float) -> None:
        self.num_pulls_per_arm[arm_id] = self.num_pulls_per_arm[arm_id] +
1
        self.Q[arm_id] = self.Q[arm_id] + (arm_reward - self.Q[arm_id]) /
self.num_pulls_per_arm[arm_id]

    def select_arm(self) -> int:
        max_q_value = max(self.Q.values())
        exp_values = {id: np.exp((self.Q[id] - max_q_value) / self.tau)
for id in self.arm_ids}
        sum_exp_values = sum(exp_values.values())
        probabilities = {id: exp_value / sum_exp_values for id, exp_value
in exp_values.items()}

```

```

        selected_arm = np.random.choice(list(probabilities.keys()),
p=list(probabilities.values()))
        return selected_arm

class UCB(BasePolicy):
    def __init__(self, c, arm_ids):
        self.c = c
        self.arm_ids = arm_ids
        self.Q = {id: 0 for id in self.arm_ids}
        self.num_pulls_per_arm = {id: 0 for id in self.arm_ids}
        self.total_pulls = 0

    @property
    def name(self):
        return f'UCB c:{self.c}'

    def reset(self):
        self.Q = {id: 0 for id in self.arm_ids}
        self.num_pulls_per_arm = {id: 0 for id in self.arm_ids}
        self.total_pulls = 0

    def update_arm(self, arm_id: int, arm_reward: float) -> None:
        self.num_pulls_per_arm[arm_id] =
self.num_pulls_per_arm[arm_id] + 1
        self.total_pulls = self.total_pulls + 1
        old_q = self.Q[arm_id]
        n = self.num_pulls_per_arm[arm_id]
        new_q = old_q + (arm_reward - old_q) / n
        self.Q[arm_id] = new_q

    def select_arm(self) -> int:
        exploration_bonus = self.c * np.sqrt(np.log(self.total_pulls +
1) / np.maximum(1, list(self.num_pulls_per_arm.values())))
        ucb_values = [self.Q[arm_id] + exploration_bonus[arm_id] for
arm_id in self.arm_ids]
        chosen_arm = max(self.arm_ids, key=lambda x: ucb_values[x])
        return chosen_arm

```

## Trainer

```

def train(env, policy: BasePolicy, timesteps):
    policy_reward = np.zeros((timesteps,))
    for t in range(timesteps):
        arm_id = policy.select_arm()
        reward = env.step(arm_id)
        policy.update_arm(arm_id, reward)
        policy_reward[t] = reward
    return policy_reward

def avg_over_runs(env, policy: BasePolicy, timesteps, num_runs):

```

```

_, expected_max_reward = env.get_best_arm_and_expected_reward()
policy_reward_each_run = np.zeros((num_runs, timesteps))

for run in range(num_runs):
    policy.reset()
    policy_reward = train(env, policy, timesteps)
    policy_reward_each_run[run, :] = policy_reward

avg_policy_rewards = np.mean(policy_reward_each_run, axis=0)
total_policy_regret = expected_max_reward * timesteps -
np.sum(avg_policy_rewards)

return avg_policy_rewards, total_policy_regret

def plot_reward_curve_and_print_regret(env, policies, timesteps=200,
num_runs=500):
    fig, ax = plt.subplots(1, 1, sharex=False, sharey=False,
figsize=(10, 6))
    for policy in policies:
        avg_policy_rewards, total_policy_regret = avg_over_runs(env,
policy, timesteps, num_runs)
        print('regret for {}: {:.3f}'.format(policy.name,
total_policy_regret))
        ax.plot(np.arange(timesteps), avg_policy_rewards, '-',
label=policy.name)

_, expected_max_reward = env.get_best_arm_and_expected_reward()
ax.plot(np.arange(timesteps), [expected_max_reward]*timesteps, 'g-')

avg_arm_reward = env.get_avg_arm_reward()
ax.plot(np.arange(timesteps), [avg_arm_reward]*timesteps, 'r-')

plt.legend(loc='lower right')
plt.show()

```

## Experiments

```


seed = 42
np.random.seed(seed)

num_arms = 5
mean_reward_range = (-25, 25)
std = 2.0

env = Env(num_arms, mean_reward_range, std)

env.plot_arms_reward_distribution()

```



```
best_arm, max_mean_reward = env.get_best_arm_and_expected_reward()
print(best_arm, max_mean_reward)
```

```
1 22.53571532049581
```

```
print(env.get_avg_arm_reward())
```


```
3.119254917081568
```

Please explore following values:

- Epsilon greedy: [0.001, 0.01, 0.5, 0.9]
- Softmax: [0.001, 1.0, 5.0, 50.0]

```
random_policy = RandomPolicy(env.arm_ids)
plot_reward_curve_and_print_regret(env, [random_policy],
timesteps=200, num_runs=500)
```


```
regret for random: 3871.625
```



```
explore_epgreedy_epsilons = [0.001, 0.01, 0.5, 0.9]
epgreedy_policies = [EpGreedyPolicy(ep, env.arm_ids) for ep in
explore_epgreedy_epsilons]
plot_reward_curve_and_print_regret(env, epgreedy_policies,
timesteps=200, num_runs=500)
```

```
regret for ep-greedy ep:0.001: 33.248
regret for ep-greedy ep:0.01: 85.360
regret for ep-greedy ep:0.5: 1992.935
regret for ep-greedy ep:0.9: 3497.860
```





```


explore_softmax_taus = [0.001, 1.0, 5.0, 50.0]
softmax_policies = [SoftmaxPolicy(tau, env.arm_ids) for tau in
explore_softmax_taus]
plot_reward_curve_and_print_regret(env, softmax_policies,
timesteps=200, num_runs=500)

```

```


regret for softmax tau:0.001: 1833.224
regret for softmax tau:1.0: 1368.913
regret for softmax tau:5.0: 400.365
regret for softmax tau:50.0: 3149.831

```



```
plot_reward_curve_and_print_regret(env, [UCB(4, env.arm_ids)],
timesteps=200, num_runs=500)
```

regret for UCB c:4: 31.505



Optional: Please explore different values of epsilon, tau and verify how does the behaviour changes.